

```

import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import cv2
import os
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img,
img_to_array
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Input,
Dropout, Conv2D, BatchNormalization, Activation, Add, MaxPooling2D
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau,
ModelCheckpoint
from tensorflow.keras.optimizers import Adam
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
import seaborn as sns
import pandas as pd
from PIL import Image
import warnings

warnings.filterwarnings('ignore', category=UserWarning, module='keras')

# Suppress CUDA warnings
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'

# Print TensorFlow version and GPU availability
print("TensorFlow Version:", tf.__version__)
print("Num GPUs Available:", len(tf.config.list_physical_devices('GPU')))

# Define paths to dataset directories
train_dir = '/kaggle/input/chest-xray-pneumonia/chest_xray/train'
test_dir = '/kaggle/input/chest-xray-pneumonia/chest_xray/test'

# Verify dataset directories
if not os.path.exists(train_dir):
    raise FileNotFoundError(f"Training directory not found: {train_dir}")
if not os.path.exists(test_dir):
    raise FileNotFoundError(f"Test directory not found: {test_dir}")

# Custom ImageDataGenerator
def create_datagen():
    return ImageDataGenerator(
        rescale=1./255,
        rotation_range=20,
        width_shift_range=0.2,

```

```

        height_shift_range=0.2,
        zoom_range=0.2,
        shear_range=0.2,
        brightness_range=[0.8, 1.2],
        horizontal_flip=True,
        fill_mode='nearest'
    )

# Oversample Normal class
def oversample_normal_images(train_dir):
    normal_dir = os.path.join(train_dir, 'NORMAL')
    pneumonia_dir = os.path.join(train_dir, 'PNEUMONIA')
    normal_files = [os.path.join(normal_dir, f) for f in os.listdir(normal_dir)
if f.endswith('.jpeg')]
    pneumonia_files = [os.path.join(pneumonia_dir, f) for f in
os.listdir(pneumonia_dir) if f.endswith('.jpeg')]

    # Oversample Normal to match Pneumonia count (~3875)
    normal_oversampled = normal_files * (len(pneumonia_files) //
len(normal_files)) + normal_files[:len(pneumonia_files) % len(normal_files)]
    files = normal_oversampled + pneumonia_files
    labels = [0] * len(normal_oversampled) + [1] * len(pneumonia_files)
    return files, labels

# Load and balance training data
train_files, train_labels = oversample_normal_images(train_dir)
train_df = pd.DataFrame({'filename': train_files, 'class': train_labels})
train_df['class'] = train_df['class'].map({0: 'NORMAL', 1: 'PNEUMONIA'})

# Split into training and validation sets (70% train, 30% validation)
train_df, val_df = train_test_split(train_df, test_size=0.3,
stratify=train_df['class'], random_state=42)

# Print DataFrame info
print("Training DataFrame head:\n", train_df.head())
print("Validation DataFrame head:\n", val_df.head())
print("Training class counts:\n", train_df['class'].value_counts())
print("Validation class counts:\n", val_df['class'].value_counts())

# Data generators
train_datagen = create_datagen()
val_datagen = create_datagen()
test_datagen = ImageDataGenerator(rescale=1./255)

batch_size = 16

```

```

training_data = train_datagen.flow_from_dataframe(
    train_df,
    x_col='filename',
    y_col='class',
    target_size=(224, 224),
    color_mode='grayscale',
    class_mode='binary',
    batch_size=batch_size,
    shuffle=True,
    seed=42
)

validation_data = val_datagen.flow_from_dataframe(
    val_df,
    x_col='filename',
    y_col='class',
    target_size=(224, 224),
    color_mode='grayscale',
    class_mode='binary',
    batch_size=batch_size,
    shuffle=True, # Ensure shuffling for balanced batches
    seed=42
)

testing_data = test_datagen.flow_from_directory(
    test_dir,
    target_size=(224, 224),
    color_mode='grayscale',
    class_mode='binary',
    batch_size=batch_size,
    shuffle=False
)

# Print class distributions
print("Training class distribution:", np.bincount(training_data.classes))
print("Validation class distribution:", np.bincount(validation_data.classes))
print("Class indices:", training_data.class_indices)

# Debug validation batch
images, labels = next(validation_data)
print("Validation batch images shape:", images.shape) # Should be (16, 224, 224, 1)
print("Validation batch labels shape:", labels.shape) # Should be (16,)
print("Validation batch labels:", labels) # Should contain 0s and 1s

```

```

# Visualize sample training and validation images
def visualize_samples(df, title, num_samples=5):
    normal_files = df[df['class'] == 'NORMAL']['filename'].sample(num_samples,
random_state=42).tolist()
    pneumonia_files = df[df['class'] ==
'PNEUMONIA']['filename'].sample(num_samples, random_state=42).tolist()

    plt.figure(figsize=(15, 6))
    for i, file in enumerate(normal_files + pneumonia_files):
        img = load_img(file, target_size=(224, 224), color_mode='grayscale')
        img_array = img_to_array(img) / 255.0
        plt.subplot(2, num_samples, i + 1)
        plt.imshow(img_array.squeeze(), cmap='gray')
        plt.title('NORMAL' if i < num_samples else 'PNEUMONIA')
        plt.axis('off')
    plt.suptitle(title)
    plt.tight_layout()
    plt.savefig(f'{title.lower().replace(" ", "_")}.png')
    plt.show()

visualize_samples(train_df, "Training Samples")
visualize_samples(val_df, "Validation Samples")

# Custom CNN with skip connections
inputs = Input(shape=(224, 224, 1))

# Block 1
x = Conv2D(32, (3, 3), padding='same')(inputs)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Conv2D(32, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Block 2 with skip connection
identity = x
x = Conv2D(64, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Conv2D(64, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
identity = Conv2D(64, (1, 1), padding='same')(identity)
x = Add()([x, identity])

```

```

x = Activation('relu')(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Block 3 with skip connection
identity = x
x = Conv2D(128, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Conv2D(128, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
identity = Conv2D(128, (1, 1), padding='same')(identity)
x = Add()(x, identity)
x = Activation('relu')(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Block 4 with skip connection
identity = x
x = Conv2D(256, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Conv2D(256, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
identity = Conv2D(256, (1, 1), padding='same')(identity)
x = Add()(x, identity)
x = Activation('relu')(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Block 5 with skip connection
identity = x
x = Conv2D(512, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Conv2D(512, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
identity = Conv2D(512, (1, 1), padding='same')(identity)
x = Add()(x, identity)
x = Activation('relu')(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Block 6 with skip connection
identity = x
x = Conv2D(1024, (3, 3), padding='same')(x)
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = Conv2D(1024, (3, 3), padding='same')(x)

```

```

x = BatchNormalization()(x)
identity = Conv2D(1024, (1, 1), padding='same')(identity)
x = Add()([x, identity])
x = Activation('relu')(x)
x = MaxPooling2D(pool_size=(2, 2))(x)

# Classification head
x = GlobalAveragePooling2D()(x)
x = Dense(256, activation='relu',
kernel_regularizer=tf.keras.regularizers.l2(0.02))(x)
x = Dropout(0.6)(x)
outputs = Dense(1, activation='sigmoid')(x)

model = Model(inputs, outputs)

# Compile model
model.compile(
    loss='binary_crossentropy',
    optimizer=Adam(learning_rate=5e-5),
    metrics=['accuracy', tf.keras.metrics.Precision(), tf.keras.metrics.Recall()]
)

# Class weights
class_weight_dict = {0: 1.5, 1: 0.5} # Prioritize Normal class

# Callbacks
early_stopping = EarlyStopping(monitor='val_loss', patience=20,
restore_best_weights=True)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=7,
min_lr=1e-6)
checkpoint = ModelCheckpoint('best_model.keras', monitor='val_accuracy',
save_best_only=True, mode='max')

# Train model
history = model.fit(
    training_data,
    epochs=70,
    validation_data=validation_data,
    class_weight=class_weight_dict,
    callbacks=[early_stopping, reduce_lr, checkpoint]
)

# Load best model
model.load_weights('best_model.keras')

```

```

# Evaluate on validation data
val_loss, val_accuracy, val_precision, val_recall =
model.evaluate(validation_data)
print(f"Validation Accuracy: {val_accuracy:.4f}, Loss: {val_loss:.4f}, Precision:
{val_precision:.4f}, Recall: {val_recall:.4f}")

# Evaluate on test data
test_loss, test_accuracy, test_precision, test_recall =
model.evaluate(testing_data)
print(f"Test Accuracy: {test_accuracy:.4f}, Precision: {test_precision:.4f},
Recall: {test_recall:.4f}")

# Predict on test data
predictions = model.predict(testing_data)
predicted_classes = (predictions > 0.5).astype(int).flatten()
true_classes = testing_data.classes

# Confusion Matrix
cm = confusion_matrix(true_classes, predicted_classes)
print("Confusion Matrix:")
print(cm)

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.savefig('confusion_matrix.png')
plt.show()

# Accuracy and Loss Plots
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')

```

```

plt.title('Training and Validation Loss')
plt.legend()
plt.tight_layout()
plt.savefig('accuracy_loss_curves.png')
plt.show()

# Classification Report
print("\nClassification Report for Test Set:")
print(classification_report(true_classes, predicted_classes,
target_names=['NORMAL', 'PNEUMONIA']))

# Data Quality Check: Visualize sample images
def visualize_samples(directory, num_samples=5):
    normal_dir = os.path.join(directory, 'NORMAL')
    pneumonia_dir = os.path.join(directory, 'PNEUMONIA')
    normal_files = [os.path.join(normal_dir, f) for f in os.listdir(normal_dir)
if f.endswith('.jpeg')][:num_samples]
    pneumonia_files = [os.path.join(pneumonia_dir, f) for f in
os.listdir(pneumonia_dir) if f.endswith('.jpeg')][:num_samples]

    plt.figure(figsize=(15, 6))
    for i, file in enumerate(normal_files + pneumonia_files):
        img = load_img(file, target_size=(224, 224), color_mode='grayscale')
        img_array = img_to_array(img) / 255.0
        plt.subplot(2, num_samples, i + 1)
        plt.imshow(img_array.squeeze(), cmap='gray')
        plt.title('NORMAL' if i < num_samples else 'PNEUMONIA')
        plt.axis('off')
    plt.tight_layout()
    plt.savefig('sample_images.png')
    plt.show()

```

```
visualize_samples(train_dir)
```

	filename	class
5197	/kaggle/input/chest-xray-pneumonia/chest_xray/...	PNEUMONIA
2399	/kaggle/input/chest-xray-pneumonia/chest_xray/...	NORMAL
2813	/kaggle/input/chest-xray-pneumonia/chest_xray/...	NORMAL
6588	/kaggle/input/chest-xray-pneumonia/chest_xray/...	PNEUMONIA
2920	/kaggle/input/chest-xray-pneumonia/chest_xray/...	NORMAL

Validation DataFrame head:

	filename	class
1346	/kaggle/input/chest-xray-pneumonia/chest_xray/...	NORMAL
1753	/kaggle/input/chest-xray-pneumonia/chest_xray/...	NORMAL
5878	/kaggle/input/chest-xray-pneumonia/chest_xray/...	PNEUMONIA
47	/kaggle/input/chest-xray-pneumonia/chest_xray/...	NORMAL
5699	/kaggle/input/chest-xray-pneumonia/chest_xray/...	PNEUMONIA

Training class counts:

class

PNEUMONIA 2713

NORMAL 2712

Name: count, dtype: int64

Validation class counts:

class

NORMAL 1163

PNEUMONIA 1162

Name: count, dtype: int64

Found 5425 validated image filenames belonging to 2 classes.

Found 2325 validated image filenames belonging to 2 classes.

Found 624 images belonging to 2 classes.

Training class distribution: [2712 2713]


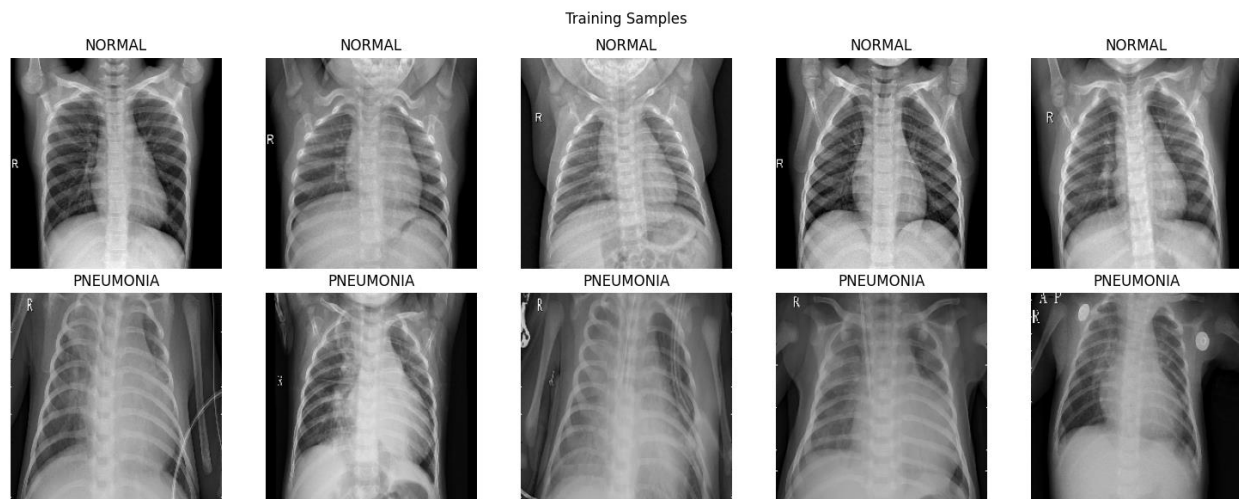
Validation class distribution: [1163 1162]

Class indices: {'NORMAL': 0, 'PNEUMONIA': 1}

Validation batch images shape: (16, 224, 224, 1)

Validation batch labels shape: (16,)

Validation batch labels: [1. 0. 0. 1. 0. 1. 0. 0. 0. 0. 1. 0. 1. 0. 1.]



I0000 00:00:1745694223.003736 31 gpu_device.cc:2022] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 13942 MB memory: -> device: 0, name: Tesla T4, pci
bus id: 0000:00:04.0, compute capability: 7.5

I0000 00:00:1745694223.004405 31 gpu_device.cc:2022] Created device
/job:localhost/replica:0/task:0/device:GPU:1 with 13942 MB memory: -> device: 1, name: Tesla T4, pci
bus id: 0000:00:05.0, compute capability: 7.5

Epoch 1/70

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1745694239.730685 103 service.cc:148] XLA service 0x782500003e40 initialized for
platform CUDA (this does not guarantee that XLA will be used). Devices:

I0000 00:00:1745694239.731681 103 service.cc:156] StreamExecutor device (0): Tesla T4, Compute
Capability 7.5

I0000 00:00:1745694239.731707 103 service.cc:156] StreamExecutor device (1): Tesla T4, Compute
Capability 7.5

I0000 00:00:1745694241.085003 103 cuda_dnn.cc:529] Loaded cuDNN version 90300

[1m 1/340[0m [37m[0m [1m3:05:22[0m

33s/step - accuracy: 0.5625 - loss: 9.8745 - precision: 0.5333 - recall: 1.0000I0000

00:00:1745694257.622962 103 device_compiler.h:188] Compiled cluster using XLA! This line is logged
at most once for the lifetime of the process.

[1m340/340[0m [32m[0m [37m[0m

[1m142s[0m 322ms/step - accuracy: 0.7961 - loss: 7.9182 - precision: 0.8840 - recall: 0.6825 -
val_accuracy: 0.4998 - val_loss: 10.9099 - val_precision: 0.4998 - val_recall: 1.0000 - learning_rate:
5.0000e-05

Epoch 2/70

[1m340/340[0m [32m[0m [37m[0m

[1m94s[0m 273ms/step - accuracy: 0.8848 - loss: 5.3298 - precision: 0.9639 - recall: 0.7965 -
val_accuracy: 0.8877 - val_loss: 4.1780 - val_precision: 0.9395 - val_recall: 0.8287 - learning_rate:
5.0000e-05

Epoch 3/70

[1m340/340[0m [32m[0m [37m[0m

[1m91s[0m 265ms/step - accuracy: 0.8996 - loss: 3.7124 - precision: 0.9708 - recall: 0.8258 -
val_accuracy: 0.9239 - val_loss: 2.8855 - val_precision: 0.9833 - val_recall: 0.8623 - learning_rate:
5.0000e-05

Epoch 4/70

[1m340/340[0m [32m[0m [37m[0m

[1m89s[0m 259ms/step - accuracy: 0.9086 - loss: 2.6471 - precision: 0.9781 - recall: 0.8387 -

val_accuracy: 0.8951 - val_loss: 2.2914 - val_precision: 0.9935 - val_recall: 0.7952 - learning_rate: 5.0000e-05

Epoch 5/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 264ms/step - accuracy: 0.9240 - loss: 1.9265 - precision: 0.9777 - recall: 0.8683 -
val_accuracy: 0.9071 - val_loss: 1.6356 - val_precision: 0.9958 - val_recall: 0.8176 - learning_rate:
5.0000e-05

Epoch 6/70

1m340/340 0m 32m 0m 37m 0m
1m92s 0m 267ms/step - accuracy: 0.9226 - loss: 1.4482 - precision: 0.9723 - recall: 0.8692 -
val_accuracy: 0.9527 - val_loss: 1.1835 - val_precision: 0.9622 - val_recall: 0.9423 - learning_rate:
5.0000e-05

Epoch 7/70

1m340/340 0m 32m 0m 37m 0m
1m92s 0m 267ms/step - accuracy: 0.9241 - loss: 1.1178 - precision: 0.9736 - recall: 0.8711 -
val_accuracy: 0.9531 - val_loss: 0.9316 - val_precision: 0.9451 - val_recall: 0.9621 - learning_rate:
5.0000e-05

Epoch 8/70

1m340/340 0m 32m 0m 37m 0m
1m92s 0m 265ms/step - accuracy: 0.9375 - loss: 0.8725 - precision: 0.9805 - recall: 0.8921 -
val_accuracy: 0.9488 - val_loss: 0.7335 - val_precision: 0.9888 - val_recall: 0.9079 - learning_rate:
5.0000e-05

Epoch 9/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 261ms/step - accuracy: 0.9441 - loss: 0.6738 - precision: 0.9829 - recall: 0.9028 -
val_accuracy: 0.9282 - val_loss: 0.6361 - val_precision: 0.9863 - val_recall: 0.8683 - learning_rate:
5.0000e-05

Epoch 10/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 260ms/step - accuracy: 0.9311 - loss: 0.5565 - precision: 0.9791 - recall: 0.8810 -
val_accuracy: 0.9475 - val_loss: 0.4940 - val_precision: 0.9943 - val_recall: 0.9002 - learning_rate:
5.0000e-05

Epoch 11/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 261ms/step - accuracy: 0.9514 - loss: 0.4326 - precision: 0.9851 - recall: 0.9170 -

val_accuracy: 0.9484 - val_loss: 0.4341 - val_precision: 0.9264 - val_recall: 0.9742 - learning_rate: 5.0000e-05

Epoch 12/70

1m340/340 [0m 32m] [0m 37m] [0m 1m90s] 260ms/step - accuracy: 0.9353 - loss: 0.3787 - precision: 0.9765 - recall: 0.8935 - val_accuracy: 0.9243 - val_loss: 0.4058 - val_precision: 0.9940 - val_recall: 0.8537 - learning_rate: 5.0000e-05

Epoch 13/70

1m340/340 [0m 32m] [0m 37m] [0m 1m91s] 264ms/step - accuracy: 0.9516 - loss: 0.3025 - precision: 0.9842 - recall: 0.9183 - val_accuracy: 0.9583 - val_loss: 0.2922 - val_precision: 0.9917 - val_recall: 0.9243 - learning_rate: 5.0000e-05

Epoch 14/70

1m340/340 [0m 32m] [0m 37m] [0m 1m90s] 260ms/step - accuracy: 0.9529 - loss: 0.2498 - precision: 0.9860 - recall: 0.9202 - val_accuracy: 0.9437 - val_loss: 0.3159 - val_precision: 0.9981 - val_recall: 0.8890 - learning_rate: 5.0000e-05

Epoch 15/70

1m340/340 [0m 32m] [0m 37m] [0m 1m90s] 260ms/step - accuracy: 0.9500 - loss: 0.2262 - precision: 0.9868 - recall: 0.9130 - val_accuracy: 0.9510 - val_loss: 0.2408 - val_precision: 0.9962 - val_recall: 0.9053 - learning_rate: 5.0000e-05

Epoch 16/70

1m340/340 [0m 32m] [0m 37m] [0m 1m89s] 257ms/step - accuracy: 0.9529 - loss: 0.1864 - precision: 0.9859 - recall: 0.9187 - val_accuracy: 0.9153 - val_loss: 0.3822 - val_precision: 0.9990 - val_recall: 0.8313 - learning_rate: 5.0000e-05

Epoch 17/70

1m340/340 [0m 32m] [0m 37m] [0m 1m91s] 262ms/step - accuracy: 0.9538 - loss: 0.1652 - precision: 0.9864 - recall: 0.9225 - val_accuracy: 0.9652 - val_loss: 0.1787 - val_precision: 0.9954 - val_recall: 0.9346 - learning_rate: 5.0000e-05

Epoch 18/70

1m340/340 [0m 32m] [0m 37m] [0m 1m91s] 263ms/step - accuracy: 0.9635 - loss: 0.1419 - precision: 0.9913 - recall: 0.9358 -

val_accuracy: 0.9622 - val_loss: 0.1544 - val_precision: 0.9590 - val_recall: 0.9656 - learning_rate: 5.0000e-05

Epoch 19/70

1m340/340 [0m 32m] [0m 37m] [0m 37m]
1m93s [0m 268ms/step - accuracy: 0.9623 - loss: 0.1288 - precision: 0.9879 - recall: 0.9364 -
val_accuracy: 0.9234 - val_loss: 0.2591 - val_precision: 0.9980 - val_recall: 0.8485 - learning_rate: 5.0000e-05

Epoch 20/70

1m340/340 [0m 32m] [0m 37m] [0m 37m]
1m92s [0m 268ms/step - accuracy: 0.9527 - loss: 0.1352 - precision: 0.9831 - recall: 0.9209 -
val_accuracy: 0.9669 - val_loss: 0.1244 - val_precision: 0.9831 - val_recall: 0.9501 - learning_rate: 5.0000e-05

Epoch 21/70

1m340/340 [0m 32m] [0m 37m] [0m 37m]
1m91s [0m 264ms/step - accuracy: 0.9609 - loss: 0.1109 - precision: 0.9864 - recall: 0.9333 -
val_accuracy: 0.9553 - val_loss: 0.1574 - val_precision: 0.9962 - val_recall: 0.9139 - learning_rate: 5.0000e-05

Epoch 22/70

1m340/340 [0m 32m] [0m 37m] [0m 37m]
1m92s [0m 265ms/step - accuracy: 0.9692 - loss: 0.0903 - precision: 0.9910 - recall: 0.9464 -
val_accuracy: 0.9677 - val_loss: 0.1219 - val_precision: 0.9936 - val_recall: 0.9415 - learning_rate: 5.0000e-05

Epoch 23/70

1m340/340 [0m 32m] [0m 37m] [0m 37m]
1m91s [0m 263ms/step - accuracy: 0.9677 - loss: 0.0920 - precision: 0.9897 - recall: 0.9444 -
val_accuracy: 0.9118 - val_loss: 0.2630 - val_precision: 1.0000 - val_recall: 0.8236 - learning_rate: 5.0000e-05

Epoch 24/70

1m340/340 [0m 32m] [0m 37m] [0m 37m]
1m90s [0m 262ms/step - accuracy: 0.9660 - loss: 0.0907 - precision: 0.9876 - recall: 0.9443 -
val_accuracy: 0.9019 - val_loss: 0.3692 - val_precision: 1.0000 - val_recall: 0.8038 - learning_rate: 5.0000e-05

Epoch 25/70

1m340/340 [0m 32m] [0m 37m] [0m 37m]
1m91s [0m 263ms/step - accuracy: 0.9670 - loss: 0.0813 - precision: 0.9882 - recall: 0.9447 -

val_accuracy: 0.9570 - val_loss: 0.1244 - val_precision: 0.9935 - val_recall: 0.9200 - learning_rate: 5.0000e-05

Epoch 26/70

1m340/340 [0m 32m] [0m 37m] [0m 1m90s] 260ms/step - accuracy: 0.9672 - loss: 0.0819 - precision: 0.9886 - recall: 0.9447 - val_accuracy: 0.8761 - val_loss: 0.3150 - val_precision: 1.0000 - val_recall: 0.7522 - learning_rate: 5.0000e-05

Epoch 27/70

1m340/340 [0m 32m] [0m 37m] [0m 1m91s] 264ms/step - accuracy: 0.9709 - loss: 0.0717 - precision: 0.9913 - recall: 0.9482 - val_accuracy: 0.9789 - val_loss: 0.0769 - val_precision: 0.9912 - val_recall: 0.9664 - learning_rate: 5.0000e-05

Epoch 28/70

1m340/340 [0m 32m] [0m 37m] [0m 1m89s] 258ms/step - accuracy: 0.9697 - loss: 0.0681 - precision: 0.9908 - recall: 0.9489 - val_accuracy: 0.9695 - val_loss: 0.1005 - val_precision: 0.9573 - val_recall: 0.9828 - learning_rate: 5.0000e-05

Epoch 29/70

1m340/340 [0m 32m] [0m 37m] [0m 1m90s] 260ms/step - accuracy: 0.9599 - loss: 0.0865 - precision: 0.9870 - recall: 0.9335 - val_accuracy: 0.9690 - val_loss: 0.0960 - val_precision: 0.9572 - val_recall: 0.9819 - learning_rate: 5.0000e-05

Epoch 30/70

1m340/340 [0m 32m] [0m 37m] [0m 1m90s] 260ms/step - accuracy: 0.9720 - loss: 0.0618 - precision: 0.9877 - recall: 0.9548 - val_accuracy: 0.9535 - val_loss: 0.1234 - val_precision: 0.9203 - val_recall: 0.9931 - learning_rate: 5.0000e-05

Epoch 31/70

1m340/340 [0m 32m] [0m 37m] [0m 1m90s] 260ms/step - accuracy: 0.9701 - loss: 0.0708 - precision: 0.9891 - recall: 0.9514 - val_accuracy: 0.8443 - val_loss: 0.5374 - val_precision: 1.0000 - val_recall: 0.6885 - learning_rate: 5.0000e-05

Epoch 32/70

1m340/340 [0m 32m] [0m 37m] [0m 1m89s] 259ms/step - accuracy: 0.9735 - loss: 0.0634 - precision: 0.9918 - recall: 0.9536 -

val_accuracy: 0.9690 - val_loss: 0.0914 - val_precision: 0.9919 - val_recall: 0.9458 - learning_rate: 5.0000e-05

Epoch 33/70

1m340/340 0m 32m 0m 37m 0m
1m89s 0m 257ms/step - accuracy: 0.9702 - loss: 0.0635 - precision: 0.9881 - recall: 0.9514 -
val_accuracy: 0.9738 - val_loss: 0.0800 - val_precision: 0.9885 - val_recall: 0.9587 - learning_rate:
5.0000e-05

Epoch 34/70

1m340/340 0m 32m 0m 37m 0m
1m89s 0m 258ms/step - accuracy: 0.9764 - loss: 0.0529 - precision: 0.9942 - recall: 0.9588 -
val_accuracy: 0.9751 - val_loss: 0.0699 - val_precision: 0.9742 - val_recall: 0.9759 - learning_rate:
5.0000e-05

Epoch 35/70

1m340/340 0m 32m 0m 37m 0m
1m89s 0m 258ms/step - accuracy: 0.9725 - loss: 0.0645 - precision: 0.9896 - recall: 0.9551 -
val_accuracy: 0.9669 - val_loss: 0.1066 - val_precision: 0.9945 - val_recall: 0.9389 - learning_rate:
5.0000e-05

Epoch 36/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 262ms/step - accuracy: 0.9712 - loss: 0.0587 - precision: 0.9922 - recall: 0.9511 -
val_accuracy: 0.9480 - val_loss: 0.1546 - val_precision: 0.9990 - val_recall: 0.8967 - learning_rate:
5.0000e-05

Epoch 37/70

1m340/340 0m 32m 0m 37m 0m
1m89s 0m 259ms/step - accuracy: 0.9747 - loss: 0.0595 - precision: 0.9946 - recall: 0.9554 -
val_accuracy: 0.9561 - val_loss: 0.1385 - val_precision: 0.9953 - val_recall: 0.9165 - learning_rate:
5.0000e-05

Epoch 38/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 262ms/step - accuracy: 0.9649 - loss: 0.0687 - precision: 0.9871 - recall: 0.9411 -
val_accuracy: 0.9708 - val_loss: 0.0949 - val_precision: 0.9991 - val_recall: 0.9423 - learning_rate:
5.0000e-05

Epoch 39/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 263ms/step - accuracy: 0.9781 - loss: 0.0531 - precision: 0.9942 - recall: 0.9614 -

val_accuracy: 0.9785 - val_loss: 0.0661 - val_precision: 0.9728 - val_recall: 0.9845 - learning_rate: 5.0000e-05

Epoch 40/70

1m340/340 [0m 32m] [0m 37m] [0m 268ms/step - accuracy: 0.9765 - loss: 0.0520 - precision: 0.9895 - recall: 0.9631 - val_accuracy: 0.9703 - val_loss: 0.0886 - val_precision: 0.9513 - val_recall: 0.9914 - learning_rate: 5.0000e-05

Epoch 41/70

1m340/340 [0m 32m] [0m 37m] [0m 275ms/step - accuracy: 0.9803 - loss: 0.0484 - precision: 0.9910 - recall: 0.9699 - val_accuracy: 0.9858 - val_loss: 0.0493 - val_precision: 0.9904 - val_recall: 0.9811 - learning_rate: 5.0000e-05

Epoch 42/70

1m340/340 [0m 32m] [0m 37m] [0m 270ms/step - accuracy: 0.9743 - loss: 0.0545 - precision: 0.9908 - recall: 0.9582 - val_accuracy: 0.9802 - val_loss: 0.0636 - val_precision: 0.9938 - val_recall: 0.9664 - learning_rate: 5.0000e-05

Epoch 43/70

1m340/340 [0m 32m] [0m 37m] [0m 260ms/step - accuracy: 0.9814 - loss: 0.0424 - precision: 0.9954 - recall: 0.9667 - val_accuracy: 0.9187 - val_loss: 0.2555 - val_precision: 0.9980 - val_recall: 0.8391 - learning_rate: 5.0000e-05

Epoch 44/70

1m340/340 [0m 32m] [0m 37m] [0m 259ms/step - accuracy: 0.9794 - loss: 0.0502 - precision: 0.9928 - recall: 0.9658 - val_accuracy: 0.9832 - val_loss: 0.0595 - val_precision: 0.9930 - val_recall: 0.9733 - learning_rate: 5.0000e-05

Epoch 45/70

1m340/340 [0m 32m] [0m 37m] [0m 261ms/step - accuracy: 0.9824 - loss: 0.0462 - precision: 0.9941 - recall: 0.9700 - val_accuracy: 0.9475 - val_loss: 0.1453 - val_precision: 0.9187 - val_recall: 0.9819 - learning_rate: 5.0000e-05

Epoch 46/70

1m340/340 [0m 32m] [0m 37m] [0m 260ms/step - accuracy: 0.9793 - loss: 0.0506 - precision: 0.9934 - recall: 0.9655 -

val_accuracy: 0.9772 - val_loss: 0.0645 - val_precision: 0.9826 - val_recall: 0.9716 - learning_rate: 5.0000e-05

Epoch 47/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 260ms/step - accuracy: 0.9815 - loss: 0.0435 - precision: 0.9933 - recall: 0.9697 -
val_accuracy: 0.9720 - val_loss: 0.1000 - val_precision: 0.9973 - val_recall: 0.9466 - learning_rate:
5.0000e-05

Epoch 48/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 264ms/step - accuracy: 0.9786 - loss: 0.0539 - precision: 0.9907 - recall: 0.9663 -
val_accuracy: 0.9716 - val_loss: 0.0810 - val_precision: 0.9832 - val_recall: 0.9596 - learning_rate:
5.0000e-05

Epoch 49/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 265ms/step - accuracy: 0.9816 - loss: 0.0493 - precision: 0.9956 - recall: 0.9670 -
val_accuracy: 0.9871 - val_loss: 0.0479 - val_precision: 0.9991 - val_recall: 0.9750 - learning_rate:
1.0000e-05

Epoch 50/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 262ms/step - accuracy: 0.9885 - loss: 0.0307 - precision: 0.9943 - recall: 0.9822 -
val_accuracy: 0.9862 - val_loss: 0.0486 - val_precision: 0.9829 - val_recall: 0.9897 - learning_rate:
1.0000e-05

Epoch 51/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 263ms/step - accuracy: 0.9891 - loss: 0.0330 - precision: 0.9959 - recall: 0.9819 -
val_accuracy: 0.9819 - val_loss: 0.0648 - val_precision: 0.9982 - val_recall: 0.9656 - learning_rate:
1.0000e-05

Epoch 52/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 262ms/step - accuracy: 0.9847 - loss: 0.0337 - precision: 0.9969 - recall: 0.9723 -
val_accuracy: 0.9871 - val_loss: 0.0414 - val_precision: 0.9871 - val_recall: 0.9871 - learning_rate:
1.0000e-05

Epoch 53/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 260ms/step - accuracy: 0.9904 - loss: 0.0271 - precision: 0.9971 - recall: 0.9835 -

val_accuracy: 0.9871 - val_loss: 0.0512 - val_precision: 0.9982 - val_recall: 0.9759 - learning_rate: 1.0000e-05

Epoch 54/70

1m340/340 0m 32m 0m 37m 0m
1m92s 0m 266ms/step - accuracy: 0.9894 - loss: 0.0294 - precision: 0.9967 - recall: 0.9818 -
val_accuracy: 0.9897 - val_loss: 0.0432 - val_precision: 0.9965 - val_recall: 0.9828 - learning_rate:
1.0000e-05

Epoch 55/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 264ms/step - accuracy: 0.9915 - loss: 0.0269 - precision: 0.9981 - recall: 0.9845 -
val_accuracy: 0.9892 - val_loss: 0.0418 - val_precision: 0.9965 - val_recall: 0.9819 - learning_rate:
1.0000e-05

Epoch 56/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 263ms/step - accuracy: 0.9910 - loss: 0.0242 - precision: 0.9979 - recall: 0.9840 -
val_accuracy: 0.9832 - val_loss: 0.0544 - val_precision: 0.9973 - val_recall: 0.9690 - learning_rate:
1.0000e-05

Epoch 57/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 262ms/step - accuracy: 0.9913 - loss: 0.0245 - precision: 0.9983 - recall: 0.9842 -
val_accuracy: 0.9871 - val_loss: 0.0443 - val_precision: 0.9905 - val_recall: 0.9836 - learning_rate:
1.0000e-05

Epoch 58/70

1m340/340 0m 32m 0m 37m 0m
1m92s 0m 267ms/step - accuracy: 0.9926 - loss: 0.0229 - precision: 0.9975 - recall: 0.9875 -
val_accuracy: 0.9892 - val_loss: 0.0396 - val_precision: 0.9880 - val_recall: 0.9905 - learning_rate:
1.0000e-05

Epoch 59/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 261ms/step - accuracy: 0.9926 - loss: 0.0230 - precision: 0.9982 - recall: 0.9869 -
val_accuracy: 0.9875 - val_loss: 0.0509 - val_precision: 0.9991 - val_recall: 0.9759 - learning_rate:
1.0000e-05

Epoch 60/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 261ms/step - accuracy: 0.9887 - loss: 0.0276 - precision: 0.9972 - recall: 0.9793 -

val_accuracy: 0.9875 - val_loss: 0.0415 - val_precision: 0.9991 - val_recall: 0.9759 - learning_rate: 1.0000e-05

Epoch 61/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 263ms/step - accuracy: 0.9869 - loss: 0.0301 - precision: 0.9936 - recall: 0.9798 -
val_accuracy: 0.9862 - val_loss: 0.0474 - val_precision: 0.9965 - val_recall: 0.9759 - learning_rate:
1.0000e-05

Epoch 62/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 261ms/step - accuracy: 0.9898 - loss: 0.0257 - precision: 0.9969 - recall: 0.9833 -
val_accuracy: 0.9871 - val_loss: 0.0358 - val_precision: 0.9974 - val_recall: 0.9768 - learning_rate:
1.0000e-05

Epoch 63/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 261ms/step - accuracy: 0.9891 - loss: 0.0262 - precision: 0.9948 - recall: 0.9834 -
val_accuracy: 0.9824 - val_loss: 0.0616 - val_precision: 1.0000 - val_recall: 0.9647 - learning_rate:
1.0000e-05

Epoch 64/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 261ms/step - accuracy: 0.9881 - loss: 0.0267 - precision: 0.9971 - recall: 0.9787 -
val_accuracy: 0.9888 - val_loss: 0.0386 - val_precision: 0.9939 - val_recall: 0.9836 - learning_rate:
1.0000e-05

Epoch 65/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 262ms/step - accuracy: 0.9874 - loss: 0.0306 - precision: 0.9942 - recall: 0.9804 -
val_accuracy: 0.9867 - val_loss: 0.0393 - val_precision: 0.9991 - val_recall: 0.9742 - learning_rate:
1.0000e-05

Epoch 66/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 263ms/step - accuracy: 0.9926 - loss: 0.0197 - precision: 0.9983 - recall: 0.9869 -
val_accuracy: 0.9824 - val_loss: 0.0532 - val_precision: 0.9991 - val_recall: 0.9656 - learning_rate:
1.0000e-05

Epoch 67/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 263ms/step - accuracy: 0.9922 - loss: 0.0204 - precision: 0.9992 - recall: 0.9850 -

val_accuracy: 0.9905 - val_loss: 0.0365 - val_precision: 0.9931 - val_recall: 0.9880 - learning_rate: 1.0000e-05

Epoch 68/70

1m340/340 0m 32m 0m 37m 0m
1m91s 0m 263ms/step - accuracy: 0.9903 - loss: 0.0231 - precision: 0.9958 - recall: 0.9847 -
val_accuracy: 0.9824 - val_loss: 0.0528 - val_precision: 0.9754 - val_recall: 0.9897 - learning_rate:
1.0000e-05

Epoch 69/70

1m340/340 0m 32m 0m 37m 0m
1m92s 0m 267ms/step - accuracy: 0.9917 - loss: 0.0240 - precision: 0.9983 - recall: 0.9852 -
val_accuracy: 0.9914 - val_loss: 0.0332 - val_precision: 0.9974 - val_recall: 0.9854 - learning_rate:
1.0000e-05

Epoch 70/70

1m340/340 0m 32m 0m 37m 0m
1m90s 0m 262ms/step - accuracy: 0.9904 - loss: 0.0240 - precision: 0.9972 - recall: 0.9835 -
val_accuracy: 0.9837 - val_loss: 0.0601 - val_precision: 0.9973 - val_recall: 0.9699 - learning_rate:
1.0000e-05

1m146/146 0m 32m 0m 37m 0m
1m27s 0m 182ms/step - accuracy: 0.9891 - loss: 0.0398 - precision: 0.9973 - recall: 0.9810

Validation Accuracy: 0.9884, Loss: 0.0361, Precision: 0.9939, Recall: 0.9828

1m39/39 0m 32m 0m 37m 0m 1m6s 0m
162ms/step - accuracy: 0.9435 - loss: 0.1825 - precision: 0.5523 - recall: 0.6294




Test Accuracy: 0.9535, Precision: 0.9524, Recall: 0.9744

1m39/39 0m 32m 0m 37m 0m 1m4s 0m
73ms/step

Confusion Matrix:

[[215 19]

[10 380]]



Classification Report for Test Set:

	precision	recall	f1-score	support
NORMAL	0.96	0.92	0.94	234
PNEUMONIA	0.95	0.97	0.96	390
accuracy		0.95		624
macro avg	0.95	0.95	0.95	624
weighted avg	0.95	0.95	0.95	624

